

MD RAKIBUL ISLAM RAIHAN

ML Systems Engineer | LLM Inference & Multimodal AI

2027 New Grad · Available July 2027 · Open to Relocation · Xi'an, China

Email: raihanrakib.143@gmail.com · Phone: +86 177 2236 2949

GitHub: github.com/theraihanrakibb · LinkedIn: linkedin.com/in/theraihanrakib · Portfolio: theraihanrakibb.github.io/Online-Portfolio

Dear Hiring Manager,

I am writing to apply for an ML Systems Engineer position focused on LLM inference and multimodal AI (2027 New Grad). I am an M.Eng. Software Engineering candidate at Northwestern Polytechnical University (Top 1%, Chinese Government Scholarship), graduating July 2027, and I am available from July 2027 and open to relocation.

At InfiX.ai I worked on production LLM serving, not demos. On 4x A800 clusters I designed a three-tier KV-cache offload pipeline over Mooncake RDMA that cut time-to-first-token 8.3x for long-context workloads, and I built prefill/decode disaggregated serving that raised stable concurrency from 192 to 256 requests. I also patched SGLang internals directly to fix production cache failures and ship hybrid-attention Qwen models, and used Nsight Systems to localize GEMM bottlenecks and apply FP8 quantization for throughput.

Alongside that systems work I am first author of NVFP4-DiT, a 4-bit audio-guided video diffusion transformer under review at IEEE TNNLS, which cuts memory 4x while preserving generation quality. Understanding both the modeling side and the serving constraints that make efficient multimodal models hard to deploy is what I would bring to your team.

I would welcome the chance to discuss how my inference and low-precision work could support your roadmap; my CV, the NVFP4-DiT manuscript, and my thesis are available on request. Thank you for your consideration.

Sincerely,

MD RAKIBUL ISLAM RAIHAN

M.Eng. Software Engineering, NWPU (Expected 2027)

MD RAKIBUL ISLAM RAIHAN

ML 系统工程师 | 大模型推理与多模态 AI

2027 届应届 · 2027 年 7 月到岗 · 可接受异地 · 中国·西安

邮箱: raihanrakib.143@gmail.com · 电话: +86 177 2236 2949 · 微信: imeddii

GitHub: github.com/theraihanrakibb · LinkedIn: linkedin.com/in/theraihanrakib · Portfolio: theraihanrakibb.github.io/Online-Portfolio



尊敬的招聘负责人：

您好！我写信申请 ML 系统工程师岗位（大模型推理与多模态方向，2027 届应届）。我是西北工业大学软件工程硕士（专业前 1%，中国政府奖学金），预计 2027 年 7 月毕业，2027 年 7 月即可到岗，并接受异地工作。

在 InfiX.ai 实习期间，我从事的是生产环境中的大模型服务，而非演示项目。在 4× A800 集群上，我基于 Mooncake RDMA 设计三级 KV-cache 卸载方案，将长上下文场景的首字延迟降低 8.3 倍；并构建 prefill/decode 分离推理架构，将稳定并发由 192 提升至 256。我还直接修改 SGLang 源码以修复生产缓存故障、支持混合注意力 Qwen 模型部署，并使用 Nsight Systems 定位 GEMM 瓶颈、以 FP8 量化提升吞吐。

在系统工作之外，我是 NVFP4-DiT 的第一作者（IEEE TNNLS 审稿中）——一个 4 比特音频引导视频扩散 Transformer，在保持生成质量的同时将显存降低 4 倍。既理解建模，也理解令高效多模态模型难以落地的服务与硬件约束，这是我能为团队带来的价值。

期待有机会进一步沟通我的推理与低精度工作如何契合贵团队的方向；我的简历、NVFP4-DiT 论文与毕业论文均可应要求提供。感谢您的阅读与考虑！

此致 敬礼

MD RAKIBUL ISLAM RAIHAN

软件工程硕士，西北工业大学（预计 2027 年毕业）