

MD RAKIBUL ISLAM RAIHAN

ML Systems Engineer | LLM Inference & Multimodal AI

2027 New Grad · Available July 2027 · Open to Relocation · Xi'an, China

Email: raihanrakib.143@gmail.com · Phone: +86 177 2236 2949

GitHub: github.com/theraihanrakibb · LinkedIn: linkedin.com/in/theraihanrakib · Portfolio: theraihanrakibb.github.io/Online-Portfolio



10 October 2026

Dear Hiring Manager,

I am writing to apply for an ML Systems Engineer position focused on LLM inference and multimodal AI, as part of the 2027 new-graduate intake. I am an M.Eng. Software Engineering candidate at Northwestern Polytechnical University, graduating in July 2027, where I rank in the top 1% of my cohort and hold the Chinese Government Scholarship. I am available from July 2027 and open to relocation.

My recent work has been production LLM serving rather than demonstrations. At InfiX.ai I designed and deployed a three-tier KV-cache offload pipeline spanning GPU HBM, CPU DRAM and SSD over Mooncake RDMA, serving Qwen3.5 and Qwen3.6-27B on 4x A800-80GB GPUs: for long-context workloads this cut time-to-first-token by 8.3x, from 55s to 6.7s, while sustaining a 92% cache hit rate. I also built prefill/decode disaggregated serving across a 2+2 GPU layout, raising stable concurrency from 192 to 256 requests, and patched SGLang's cache-management internals directly to fix production cache failures and ship hybrid-attention Qwen models.

Profiling that stack with Nsight Systems changed our priorities: GEMM accounted for 60% of step time, so the real bottleneck was compute rather than cache I/O, and we redirected effort into kernels and an FP8 KV cache that halved per-token memory from 72.9 KB to 36.5 KB. I deployed Qwen3-235B-A22B-FP8, a 235B mixture-of-experts model, on 2x H200 GPUs and benchmarked DeepSeek-V3.2 on 8x H200, reaching 91 ms TTFT and 36 ms TPOT at 32-64 concurrency and 13.16 req/s at 256.

In parallel I am first author of NVFP4-DiT, a 4-bit audio-guided video diffusion transformer under review at IEEE TNNLS, whose FP4-packed Triton kernels cut memory 4x and run 3.2x faster than FP16 at competitive quality. What I would bring to your team is both halves of that profile: the modelling decisions that determine whether a model is efficient, and the serving constraints that determine whether it can be deployed. I would welcome the chance to discuss how this work could support your roadmap; my CV, the NVFP4-DiT manuscript and my thesis are available on request. Thank you for your consideration.

Sincerely,

MD RAKIBUL ISLAM RAIHAN

M.Eng. Software Engineering, NWPU (Expected 2027)

MD RAKIBUL ISLAM RAIHAN

ML 系统工程师 | 大模型推理与多模态 AI

2027 届应届 · 2027 年 7 月到岗 · 可接受异地 · 中国·西安

邮箱: raihanrakib.143@gmail.com · 电话: +86 177 2236 2949 · 微信: imeddii

GitHub: github.com/theraihanrakibb · LinkedIn: linkedin.com/in/theraihanrakib · Portfolio: theraihanrakibb.github.io/Online-Portfolio



2026年10月10日

尊敬的招聘负责人:

您好! 我写信申请 ML 系统工程师岗位 (大模型推理与多模态方向), 为 2027 届应届毕业生。我是西北工业大学软件工程硕士, 预计 2027 年 7 月毕业, 专业排名前 1%, 并获得中国政府奖学金。我 2027 年 7 月即可到岗, 并接受异地工作安排。

我近期的工作是生产环境中的大模型服务, 而非演示项目。在 InfiX.ai, 我设计并部署了跨越 GPU HBM、CPU 内存与 SSD 的三级 KV 缓存卸载方案, 基于 Mooncake RDMA 为 4x A800-80GB 上的 Qwen3.5 与 Qwen3.6-27B 提供服务: 在长上下文场景下, 首字延迟降低 8.3 倍 (55s → 6.7s), 同时缓存命中率维持在 92%。我还构建了 2+2 GPU 的 prefill/decode 分离推理架构, 将稳定并发由 192 提升至 256, 并直接修改 SGLang 缓存管理内部模块以修复生产缓存故障、支持混合注意力 Qwen 模型部署。

使用 Nsight Systems 分析该服务栈改变了我们的优化重心: GEMM 占单步耗时 60%, 真正的瓶颈是计算而非缓存 I/O, 我们因此将精力转向算子与 FP8 KV 缓存, 将单 token 显存由 72.9 KB 减半至 36.5 KB。我在 2x H200 上部署了 235B 混合专家模型 Qwen3-235B-A22B-FP8, 并在 8x H200 上对 DeepSeek-V3.2 进行基准测试, 达到 32-64 并发下 91 ms TTFT 与 36 ms TPOT、256 并发下 13.16 req/s。

与此同时, 我是 NVFP4-DiT 的第一作者——一个 4 比特音频引导视频扩散 Transformer, 目前在 IEEE TNNLS 审稿中; 其 FP4 打包 Triton 算子在保持生成质量竞争力的同时将显存降低 4 倍, 并相较 FP16 提速 3.2 倍。我能为团队带来的是这种能力的两半: 决定模型是否高效的建模决策, 以及决定它能否真正部署的服务约束。期待有机会进一步沟通这些工作如何契合贵团队的方向; 我的简历、NVFP4-DiT 论文与毕业论文均可应要求提供。感谢您的阅读与考虑!

此致 敬礼

MD RAKIBUL ISLAM RAIHAN

软件工程硕士, 西北工业大学 (预计 2027 年毕业)