

MD RAKIBUL ISLAM RAIHAN

ML Systems Engineer | LLM Inference & Multimodal AI

2027 New Grad · Available July 2027 · Open to Relocation · Xi'an, China

Email: raihanrakib.143@gmail.com · Phone: +86 177 2236 2949

GitHub: github.com/theraihanrakibb · LinkedIn: [linkedin.com/in/theraihanrakib](https://www.linkedin.com/in/theraihanrakib) · Portfolio: theraihanrakibb.github.io/Online-Portfolio



SUMMARY

AI Systems Engineer focused on LLM inference, GPU optimization, and multimodal generative AI. At InfiX.ai, optimized production LLM serving on 4× A800/H800 GPUs, reducing TTFT 8.3× (55s→6.7s) through 3-tier KV-cache offload and improving stable concurrency 192→256 using RDMA-based prefill/decode disaggregation. First author of an IEEE TNNLS manuscript on NVFP4 4-bit audio-guided video diffusion, achieving 4× memory reduction on A800/H100. Strong in SGLang, CUDA/Triton, low-precision inference, distributed systems, and PyTorch.

EDUCATION

M.Eng. in Software Engineering, School of Software, Northwestern Polytechnical University - Sep 2024 - Jul 2027

GPA: 88/100, Top 1%, Chinese Government Scholarship, President of International Student Union (ISU)

B.Eng. in Computer Science & Technology, School of Computer Science, Northwestern Polytechnical University - Sep 2020 - Jul 2024

GPA: 85/100, Top 1%, Presidential Scholarship, Best Undergraduate Thesis Award

EXPERIENCE

AI Infrastructure Engineer Intern, InfiX.ai, Shenzhen, China, Apr 2026 - Jun 2026

- Reduced LLM time-to-first-token (TTFT) 8.3× (55s→6.7s) by designing and deploying a 3-tier KV-cache offload system across GPU HBM, CPU DRAM, and SSD using Mooncake RDMA; achieved **92%** cache hit rate for long-context workloads on 4× A800.
- Improved stable serving concurrency 33% (192→256 requests) by implementing prefill/decode disaggregation with RDMA-based KV transfer, RadixAttention prefix caching, and continuous batching.
- Modified **SGLang** cache-management and memory-allocation internals to resolve production cache failures and support hybrid-attention Qwen models.
- Profiled CUDA/Triton inference with **Nsight Systems** and optimized tensor-parallel deployment of Qwen on A800 and DeepSeek on multi-node H800 clusters.
- Diagnosed and resolved **12+** production incidents involving RDMA, etcd, storage, and distributed inference infrastructure.

Software Engineer Intern (AI Agent), Hong Kong Amram Group, Shenzhen, China, Jun 2025 - Sep 2025

- Developed production full-stack applications using React, REST APIs, and backend automation services.
- Built AI-agent workflow automation systems that improved operational efficiency by **25%**.

Electrical Software Engineer Intern, Shaanxi Longong Intelligent Technology, Xi'an, China, Feb 2025 - Apr 2025

- Built CODESYS PLC industrial-control systems with Modbus/CAN/Ethernet-IP protocols, improving motor-control performance by **15%**.

RESEARCH

Graduate Researcher, Northwestern Polytechnical University, Sep 2024 - Present

Advisor: Prof. Zhang Peng

- Developed **NVFP4-DiT**, a **4-bit audio-guided video diffusion transformer** achieving **4x memory reduction** while maintaining generation quality.
- Implemented low-precision quantization (**FP8/FP4/NVFP4**) for multimodal diffusion transformers; evaluated quality-efficiency trade-offs.
- Benchmarked NVFP4-DiT on **A800/H100 GPUs**; profiled end-to-end latency and memory with PyTorch Profiler and Nsight to guide low-precision kernel choices.
- Research on multimodal deepfake detection using visual, audio, and temporal consistency analysis.
- M.Eng. Thesis: "Detecting Deepfake Video by a Multimodal Audio-Visual Framework with Temporal Inconsistencies."

Undergraduate Researcher, Northwestern Polytechnical University, Jan 2023 - Jul 2024

Advisor: Dr. Yaxing Chen

- Designed and built a distributed confidential-query framework on Apache Spark + CKKS homomorphic encryption, achieving **95% privacy preservation** and **20% performance improvement** via optimized analytics pipelines.
- B.Eng. Thesis: "Design and Implementation of a Distributed Confidential Query Protocol for Apache Spark."

PROJECTS

AI Infra Portfolio, github.com/theraihanrakibb/ai-infra-portfolio

- 10 production-grade, Dockerized AI-infra services: LLM gateway, model serving, GPU observability, KV-cache simulator, benchmarking, quantization, tracing, and RAG.

Distributed Training Launcher, github.com/theraihanrakibb/train-launcher

- PyTorch distributed CLI for multi-node job launch, fault detection, automatic retry, and checkpoint recovery; reduces manual ops overhead for distributed training.

NVFP4-DiT (Multimodal Generation), github.com/theraihanrakibb/NVFP4-DiT

- First-author IEEE TNNLS (Q1) manuscript on 4-bit audio-guided video diffusion; 4× memory reduction while preserving generation quality on A800/H100.

TECHNICAL SKILLS

LLM Inference & Systems: SGLang, vLLM, TensorRT-LLM, KV Cache, PD Disaggregation, Mooncake/RDMA, Continuous Batching, RadixAttention

GPU & Performance: CUDA, CUDA C++, Triton, CUTLASS, FlashAttention, Nsight Systems, Nsight Compute, FP8/FP4/INT4

ML & Multimodal: PyTorch, Hugging Face Transformers, Diffusion/DiT, Multimodal Learning, OpenCV

Distributed Systems: PyTorch Distributed, DeepSpeed, Docker, Kubernetes, Linux, Prometheus, CI/CD

Programming: Python, C++, CUDA C++, Bash, SQL, Java, Data Structures & Algorithms

PUBLICATIONS & HONORS

First Author — "NVFP4-DiT: Efficient 4-Bit Audio-Guided Video Diffusion Transformers for Low-Cost Video Generation" — IEEE Transactions on Neural Networks and Learning Systems (TNNLS, Q1), manuscript under review.

Honors: Chinese Government Scholarship (2024-2027); Presidential Scholarship (2020-2024); Best Undergraduate Thesis Award (2024); Outstanding Graduate (2024), NWPU; Hackathon Winner - NWPU AI Challenge (2023).

MD RAKIBUL ISLAM RAIHAN

ML 系统工程师 | 大模型推理与多模态 AI

2027 届应届 · 2027 年 7 月到岗 · 可接受异地 · 中国·西安

邮箱: raihanrakib.143@gmail.com · 电话: +86 177 2236 2949 · 微信: imeddii

GitHub: github.com/theraihanrakibb · LinkedIn: linkedin.com/in/theraihanrakib · 作品集: theraihanrakibb.github.io



个人简介

AI 系统工程师，专注大模型推理、GPU 优化与多模态生成式 AI。在 InfiX.ai 实习期间，于 4× A800/H800 GPU 上优化生产级大模型服务，通过三级 KV-cache 卸载将 TTFT 降低 8.3× (55s→6.7s)，并借助基于 RDMA 的 prefill/decode 分离将稳定并发由 192→256 提升。作为第一作者撰写 IEEE TNNLS 论文，研究 NVFP4 4 比特音频引导视频扩散，在 A800/H100 上实现显存降低 4×。擅长 SGLang、CUDA/Triton、低精度推理、分布式系统与 PyTorch。

教育背景

软件工程硕士，软件学院，西北工业大学 (985/211) - 2024年9月 - 2027年7月

GPA: 88/100, 前 1%, 中国政府奖学金, 国际学生联合会主席 (ISU)

计算机科学与技术学士，计算机学院，西北工业大学 (985/211) - 2020年9月 - 2024年7月

GPA: 85/100, 前 1%, 校长奖学金, 最佳本科毕业论文奖

工作经历

无界智索 (深圳) 科技 (InfiX.ai), AI 基础设施工程师实习, 深圳, 2026年4月 - 2026年6月

- 在 4× A800 上基于 Mooncake RDMA 设计并实现三级 KV-cache 卸载系统 (GPU HBM → CPU DRAM → SSD)，将大模型首字延迟 (TTFT) 降低 8.3× (55s→6.7s)，长上下文工作负载缓存命中率高达 92%。
- 通过基于 RDMA 的 KV 传输实现 prefill/decode 分离，结合 RadixAttention 前缀缓存与连续批处理，将稳定服务并发提升 33% (192→256 请求)。
- 修改 SGLang 缓存管理与内存分配内部模块，修复生产缓存故障并支持混合注意力 Qwen 模型。
- 使用 Nsight Systems 分析 CUDA/Triton 推理，并优化 Qwen (A800) 与 DeepSeek (H800 多节点集群) 的张量并行部署。
- 排查并解决 12+ 起生产故障，涉及 RDMA、etcd、存储与分布式推理基础设施。

香港阿姆兰集团 (Amram Group), 软件工程师实习 (AI Agent), 深圳, 2025年6月 - 2025年9月

- 基于 React、REST API 与后端自动化服务开发生产级全栈应用。
- 开发 AI Agent 工作流自动化系统，运营效率提升 25%。

陕西龙工智能科技 (Longong), 电气软件工程师实习, 西安, 2025年2月 - 2025年4月

- 基于 CODESYS 开发 PLC 工业控制系统，实现 Modbus、CAN 与 Ethernet/IP 通信协议，电机控制性能提升 15%。

研究经历

研究生研究员, 西北工业大学, 2024年9月 - 至今

导师: 张蓬 教授

- 开发 NVFP4-DiT: 4 比特音频引导视频扩散 Transformer，显存降低 4 倍且保持生成质量。
- 实现多模态扩散 Transformer 的低精度量化 (FP8/FP4/NVFP4)，评估质量-效率权衡。
- 在 A800/H100 上对 NVFP4-DiT 进行基准测试，使用 PyTorch Profiler 与 Nsight 分析端到端延迟与显存，指导低精度算子选型。
- 基于视觉、音频与时序一致性的多模态深度伪造检测研究。
- 硕士论文: 《检测深度伪造视频: 基于多模态音视频框架与时序不一致性》。

本科研究员, 西北工业大学, 2023年1月 - 2024年7月

导师: 陈亚星 博士

- 基于 Apache Spark 与 CKKS 同态加密设计并实现分布式机密查询框架，通过优化分析管线实现 95% 隐私保护与 20% 性能提升。
- 本科论文: 《面向 Apache Spark 的分布式机密查询协议设计与实现》。

项目

AI 基础设施代码集, github.com/theraihanrakibb/ai-infra-portfolio

- 10 个生产级、Docker 化的 AI 基础设施服务: LLM 网关、模型服务、GPU 可观测性、KV-cache 模拟器、基准测试、量化、链路追踪与 RAG。

分布式训练启动器, github.com/theraihanrakibb/train-launcher

- PyTorch 分布式命令行工具，支持多节点任务启动、故障检测、自动重试与断点恢复；降低分布式训练运维成本。

NVFP4-DiT (多模态生成), github.com/theraihanrakibb/NVFP4-DiT

- 第一作者 IEEE TNNLS (Q1) 4 比特音频引导视频扩散论文；在 A800/H100 上显存降低 4× 且保持生成质量。

技术技能

大模型推理与系统: SGLang, vLLM, TensorRT-LLM, KV Cache, PD 分离, Mooncake/RDMA, 连续批处理, RadixAttention

GPU 与性能: CUDA, CUDA C++, Triton, CUTLASS, FlashAttention, Nsight Systems, Nsight Compute, FP8/FP4/INT4

机器学习与多模态: PyTorch, Hugging Face Transformers, Diffusion/DiT, 多模态学习, OpenCV

分布式系统: PyTorch Distributed, DeepSpeed, Docker, Kubernetes, Linux, Prometheus, CI/CD

编程语言: Python, C++, CUDA C++, Bash, SQL, Java, 数据结构与算法

论文与荣誉

第一作者 — 《NVFP4-DiT: 面向低成本视频生成的高效 4 比特音频引导视频扩散 Transformer》— IEEE TNNLS (Q1)，审稿中。

荣誉: 中国政府奖学金 (2024-2027)；校长奖学金 (2020-2024)；最佳本科毕业论文奖 (2024)；西北工业大学优秀毕业生 (2024)；NWPU AI 挑战赛黑客松冠军 (2023)。