

MD RAKIBUL ISLAM RAIHAN

AI Engineer — LLM Inference & Multimodal Research (2027 New Grad)

Xi'an, China · Open to Relocation · Available July 2027
Nationality: Bangladesh · Sex: Male · DOB: 2002.12.06
Languages: Bengali (Native) · English (Fluent) · Chinese (Conversational)
Work Authorization: Requires China work visa (employer-sponsored Z Visa) · Global: Sponsorship welcome
raihanrakib.143@gmail.com · +86-17722362949 · WeChat: imeddi
LinkedIn: linkedin.com/in/theraihanrakib · GitHub: github.com/theraihanrakibb
Portfolio: theraihanrakibb.github.io/Online-Portfolio



SUMMARY

2027 New Grad AI Engineer focused on LLM inference systems and multimodal research. Built and deployed LLM serving on **4× A800** clusters at **InfIX.ai** (SGLang, KV-cache offload, PD-disaggregation — **92.27% cache hit**, **8.3× faster TTFT**). **First author**, IEEE TNNLS (Q1). Seeking AI / Research Engineer roles in China, Hong Kong, and globally.

EDUCATION

M.Eng. Software Engineering, Northwestern Polytechnical University (985/211) — Sep 2024 – Jul 2027

School of Software · GPA **88/100** · **Top 1%** · Chinese Government Scholarship

B.Eng. Computer Science & Technology, Northwestern Polytechnical University (985/211) — Sep 2020 – Jul 2024

School of Computer Science · GPA **85/100** · **Top 1%** · Presidential Scholarship · Best Undergraduate Thesis Award

EXPERIENCE

AI Infrastructure Engineer Intern | InfIX.ai | Shenzhen, China | Apr 2026 – Jun 2026

- Cut TTFT 8.3× (55s → 6.7s) by building a 3-tier KV-cache offload (GPU HBM → CPU DRAM → SSD) over Mooncake RDMA on 4× A800, reaching 92.27% cache-hit for long-context models.
- Patched SGLang source across cache-management and memory-allocation modules (kvcacheio.py, hiradix_cache.py, memory_pool_host.py) to fix production cache failures and deploy hybrid-attention Qwen models.
- Built PD-disaggregated (prefill/decode) serving with RDMA KV sharing, raising stable concurrency 192 → 256 req and boosting GPU utilization via RadixAttention prefix-caching and continuous batching.
- Profiled CUDA/Triton inference with Nsight Systems, localized GEMM bottlenecks, and improved serving throughput via kernel optimization and FP8 quantization.
- Applied Little's Law to model capacity limits, proving long-context serving on 4× A800-80GB was bounded to ~1.5 RPS and guiding scale-out.
- Diagnosed and resolved 12+ production incidents across RDMA, etcd, networking, storage, and distributed inference.
- Deployed Qwen models on A800 clusters (TP2 single-node, TP4 with PD-disaggregation) and ran DeepSeek on H800 multi-node clusters (8 GPUs each), tuning tensor-parallel layouts for throughput and stability.

Software Engineer Intern (AI Agent) | Hong Kong Amram Group | Shenzhen, China | Jun 2025 – Sep 2025

- Developed production full-stack applications using React, REST APIs, and backend automation services.
- Built AI-agent workflow automation systems that improved operational efficiency by **25%**.

Electrical Software Engineer Intern | Shaanxi Longong Intelligent Technology | Xi'an, China | Feb 2025 – Apr 2025

- Developed PLC-based industrial control systems using CODESYS for robotics and industrial automation.
- Implemented Modbus, CAN, and Ethernet/IP communication protocols, improving motor-control performance by **15%**.

RESEARCH

Graduate Researcher | Northwestern Polytechnical University | Sep 2024 – Present

Advisor: Prof. Zhang Peng

- Developed **NVFP4-DiT**, a **4-bit audio-guided video diffusion transformer** achieving **4× memory reduction** while maintaining generation quality.
- Implemented low-precision quantization (**FP8/FP4/NVFP4**) for multimodal diffusion transformers; evaluated quality-efficiency trade-offs.
- Benchmarked NVFP4-DiT on **A800/H100 GPUs**; profiled end-to-end latency and memory with PyTorch Profiler and Nsight to guide low-precision kernel choices.
- Research on multimodal deepfake detection using visual, audio, and temporal consistency analysis.
- M.Eng. Thesis: "Detecting Deepfake Video by a Multimodal Audio-Visual Framework with Temporal Inconsistencies."

Undergraduate Researcher | Northwestern Polytechnical University | Jan 2023 – Jul 2024

Advisor: Dr. Yaxing Chen

- Designed and built a distributed confidential-query framework on Apache Spark + CKKS homomorphic encryption, achieving **95% privacy preservation** and **20% performance improvement** via optimized analytics pipelines.
- B.Eng. Thesis: "Design and Implementation of a Distributed Confidential Query Protocol for Apache Spark."

PROJECTS

AI Infra Portfolio | github.com/theraihanrakibb/ai-infra-portfolio

- 10 production-grade, Dockerized AI-infra projects: LLM gateway, model serving, GPU observability, training ops, MLOps registry, KV-cache simulator, benchmarking, quantization, tracing, RAG.

Distributed Training Launcher | github.com/theraihanrakibb/train-launcher

- CLI for multi-node PyTorch training with fault tolerance and automatic retry; cut manual ops overhead for distributed jobs.

C Compiler (Frontend) | github.com/theraihanrakibb/c-compiler-frontend

- Self-built C compiler (lexer/parser/codegen) using Yacc/Lex; applied systems and algorithms fundamentals.

TECHNICAL SKILLS

Languages

Python (primary), C, C++, Java, SQL, Bash

AI / ML

PyTorch, TensorFlow, Transformers, scikit-learn, NumPy; deep learning, model evaluation & metrics

LLMs & serving

LLM inference & serving (SGLang, vLLM), RAG, agents, fine-tuning (LoRA/QLoRA), quantization (FP8/INT4)

GPU & systems

CUDA, Triton, FlashAttention, multi-GPU/distributed training (DDP, FSDP), PyTorch Profiler, Nsight

Computer vision & generative AI

CNNs, Vision Transformers, diffusion models (DDPM/DiT), video generation, multimodal learning, OpenCV

Data & MLOps

Pandas, SQL, Docker, Kubernetes, Git, CI/CD (GitHub Actions), Slurm, AWS / GCP, Apache Spark, Weights & Biases / MLflow

Math & theory

linear algebra, probability & statistics, optimization, convex optimization, information theory

Fundamentals

data structures & algorithms, concurrency, operating systems, computer networks, distributed systems

PUBLICATIONS & HONORS

Publication: First Author, "NVFP4-DiT: Efficient 4-Bit Audio-Guided Video Diffusion Transformers for Low-Cost Video Generation" — IEEE Transactions on Neural Networks and Learning Systems (TNNLS, Q1), Under Review.

Honors: Chinese Government Scholarship (2024–2027); Presidential Scholarship (2020–2024); Best Undergraduate Thesis Award (2024); Outstanding Graduate (2024), NWPU; Hackathon Winner — NWPU AI Challenge (2023).